

HW 7

Evan Ott
UT EID: eao466
October 31, 2016

Laplacian smoothing

(A)

FIXME: do this

(B)

The problem is

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad \frac{1}{2} \|y - x\|_2^2 + \frac{\lambda}{2} x^\top L x = \frac{1}{2} (x^\top x - 2y^\top x + y^\top y + \lambda x^\top L x)$$

which we can find a solution for by taking the gradient w.r.t. x .

$$\begin{aligned} 0 &= \frac{1}{2} (2x - 2y + 0 + \lambda(L + L^\top)x) \\ &= \left(I + \frac{1}{2} \lambda(L + L^\top) \right) x - y \\ (I + \lambda L) \hat{x} &= y \end{aligned}$$

(C)

Gauss-Seidel

Solving $Ax = b$ in this framework amounts to splitting $A = L_* + U$ where L_* is the lower triangular matrix (including the diagonal) and U is the upper triangular matrix (excluding the diagonal).

The algorithm is then re-writing the problem iteratively as which can be solved with forward substitution.

$$L_* x^+ = b - Ux$$

See [Barrett et al. \(1994\)](#) §2.2.2 and Equation (2.6) (slightly different notation but same result).

Jacobi

Again solving $Ax = b$, we split into $A = D + R$ where D is just the diagonals and R is everything else (with 0 along the diagonal).

We then iterate

$$x^+ = D^{-1}(b - Rx)$$

What's nice here from an efficiency perspective is that D is easily invertible (so long as there are no zeros in the diagonal), and potentially linearizable depending on the underlying code implementation.

See [Barrett et al. \(1994\)](#) §2.2.1 and Equation (2.4) (slightly different notation but same result).

Conjugate Gradient

$Ax = b$ is equivalent to minimizing $\phi(x) = \frac{1}{2}x^\top Ax - b^\top x$ if A is symmetric.

$$\nabla\phi(x) = Ax - b \equiv r(x)$$

Then define $\{p_0, p_1, \dots, p_n\}$ as being conjugate w.r.t. A , such that $p_i^\top Ap_j = 0$ when $i \neq j$.

Now the algorithm proceeds as

$$\begin{aligned}x^{(k+1)} &= x^{(k)} + \alpha_k p_k \\ \alpha_k &= -\frac{r_k^\top p_k}{p_k^\top A p_k}\end{aligned}$$

So now we need to determine how to construct the p_k vectors. You could use eigenvectors, but those are in general pretty expensive to calculate. So now, let

$$\begin{aligned}p_0 &= -r_0 \\ p_k &= -r_k + \beta_k p_{k-1} \\ \beta_k &= \frac{r_{k+1}^\top r_{k+1}}{r_k^\top r_k}\end{aligned}$$

Notes from class

Next week: graph fused lasso.

Conjugate gradient is a lot more complicated, and is in fact solving a more general class of problems called Krylov subspace problems. And it is *fast* if the matrix falls into this Laplacian class of matrices (certain properties). In that case, as opposed to a standard matrix inverse, $O(n^3)$ or a sparse one, $O(n^2)$, it will actually be more like $O(n \ln n)$.

It turns out to be important to have a preconditioner. Solving $Ax = b$ is the same as solving $P^{-1}Ax = P^{-1}b$ where P is a “preconditioner.” The closer A is to P (while P is still easy to invert), the closer $P^{-1}A$ is to the identity, making the problem trivial. The current state of the art is the *algebraic multigrid* which is that $O(n \ln n)$ type solution.

Notation for the rest of the notes:

$$\begin{aligned}C_\lambda \hat{x} &= y \\ \hat{x} &= C_\lambda^{-1} y\end{aligned}$$

so \hat{x} is the smoothed/predicted y .

The question now is how to choose λ , potentially using C_p or AIC/BIC, cross-validation, etc.

The leave-one-out lemma (from [Hastie et al. \(2001\)](#)) allows us to calculate the LOOCV error. Assume $\hat{y} = Sy$ where $y, \hat{y} \in \mathbb{R}^n$ and S is a smoothing matrix (so the linear case we care about).

We need the degrees of freedom of an estimator/model (basically number of free parameters).

$$\begin{aligned}y &= X\beta + \epsilon \\ \hat{y} &= X\hat{\beta}\end{aligned}$$

has p degrees of freedom if $\beta \in \mathbb{R}^p$. We need to modify the definition to handle other cases.

Situations we need to address at a minimum:

1. Fit to p variables (as an example, OLS)

2. Choose p from $D > p$ candidate variables, then find the best fit for these p . This should have more degrees of freedom if our definition is supposed to make sense at all.

Suppose we have \hat{y} such that $\hat{y}_i = g_i(y)$. Then

$$\begin{aligned}\text{df}(\hat{y}) &= \frac{1}{\sigma^2} \sum_{i=1}^n \text{cov}(\hat{y}_i, y_i) \\ \sigma^2 &= \text{var}(y_i)\end{aligned}$$

Case: extreme overfitting: $\hat{y}_i = y_i$. Then

$$\text{df}(\hat{y}) = \frac{1}{\sigma^2} \sum_{i=1}^n \text{var}(y_i) = n$$

Case: extreme underfitting: $\hat{y}_i = \bar{y}$ Then

$$\begin{aligned}\text{df}(\hat{y}) &= \frac{1}{\sigma^2} \sum_{i=1}^n \text{cov}(\bar{y}, y_i) \\ \text{cov}(\bar{y}, y_i) &= \text{cov}\left(\frac{1}{N}y_i + \frac{1}{N} \sum_{j \neq i} y_j, y_i\right) \\ &= \text{cov}\left(\frac{1}{N}y_i, y_i\right) = \frac{1}{N}\sigma^2 \\ \text{df}(\hat{y}) &= \frac{1}{\sigma^2} \sum_{i=1}^n \frac{1}{N}\sigma^2 = 1\end{aligned}$$

Case: linear smoothers: $\hat{y} = Sy$ then $\text{df}(\hat{y}) = \text{trace}(S)$.

Graph fused lasso

Switching to an ADMM framework, we have

$$\begin{aligned}&\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad \frac{1}{2} \|y - x\|_2^2 + \lambda \|r\|_1 \\ &\text{subject to} \quad Dx - r = 0\end{aligned}$$

References

- Richard Barrett, Michael W Berry, Tony F Chan, James Demmel, June Donato, Jack Dongarra, Victor Eijkhout, Roldan Pozo, Charles Romine, and Henk Van der Vorst. *Templates for the solution of linear systems: building blocks for iterative methods*, volume 43. Siam, 1994. <http://www.netlib.org/templates/templates.pdf>.
- Trevor Hastie, Robert Tibshirani, and Jerome Friedman. The elements of statistical learning: data mining, inference and prediction. *New York: Springer-Verlag*, 1(8):371–406, 2001. <http://statweb.stanford.edu/~tibs/ElemStatLearn/>.